

Text Retrieval and Search Engines



Lecture 10 – Relevance Models



Some slides adapted from the companion website to the "Introduction to IR" book by Manning, Raghavan and Schütze & Victor Lavrenko




Motivation


- PRP: $p(R|D, q) \cong \frac{P(R|D, q)}{P(NR|D, q)} = \prod_{w \in D} \left(\frac{p_w(1-r_w)}{r_w(1-p_w)} \right)$

Motivation

▪ PRP: $p(R|D, q) \cong \frac{P(R|D, q)}{P(NR|D, q)} = \prod_{w \in D} \left(\frac{p_w(1-r_w)}{r_w(1-p_w)} \right)$



Motivation


- PRP: $p(R|D, q) \cong \frac{P(R|D, q)}{P(NR|D, q)} = \prod_{w \in D} \left(\frac{p_w(1-r_w)}{r_w(1-p_w)} \right)$


relevance model

background model
- Background model -> entire set of documents in the collection
- Relevance model -> set of known relevant documents


Motivation

- PRP: $p(R|D, q) \cong \frac{P(R|D, q)}{P(NR|D, q)} = \prod_{w \in D} \left(\frac{\overset{\text{relevance model}}{p_w(1-r_w)}}{\underset{\text{background model}}{r_w(1-p_w)}} \right)$
- Background model -> entire set of documents in the collection
- Relevance model -> set of known relevant documents



Motivation


- PRP: $p(R|D, q) \cong \frac{P(R|D, q)}{P(NR|D, q)} = \prod_{w \in D} \left(\frac{p_w(1-r_w)}{r_w(1-p_w)} \right)$
 - relevance model
 - background model
- Background model -> entire set of documents in the collection
- Relevance model -> set of known relevant documents
 - What if we have no relevant examples?




Motivation

- PRP: $p(R|D, q) \cong \frac{P(R|D, q)}{P(NR|D, q)} = \prod_{w \in D} \left(\frac{p_w(1-r_w)}{r_w(1-p_w)} \right)$
 - relevance model (points to $p_w(1-r_w)$)
 - background model (points to $r_w(1-p_w)$)
- Background model -> entire set of documents in the collection
- Relevance model -> set of known relevant documents
 - What if we have no relevant examples?
 - Assume query q is a random sample from p

Query: "hubble telescope achievements"




.080 hubble
.070 telescope
.054 earth
.002 usa



Nasa lunches a
telescope space
star earth...


Estimation: The Sampling Game

- Given query q and relevant documents $d_1, d_2, \dots, d_n$ there is a relevance language model R that generates the terms both in q and in $d_1, d_2, \dots, d_n$




Estimation: The Sampling Game

- The query and the documents are samples from an unknown relevant model R
- The sampling game
 - Relevant examples are missing
 - Let $q = q_1 \dots q_n$ - the only sample that we have from the R distribution
 - Unknown distribution P , sample 3 times, get: **hubble**, **telescope**, **achievements**
 - What's the chance we get a word **earth** if we sample one more time?




Estimation: The Sampling Game

- The query and the documents are samples from an unknown relevant model R
- The sampling game
 - Relevant examples are missing
 - Let $q = q_1 \dots q_n$ - the only sample that we have from the R distribution
 - Unknown distribution P , sample 3 times, get: **hubble**, **telescope**, **achievements**
 - What's the chance we get a word **earth** if we sample one more time?
 - $P(\text{earth}|R) \approx p(\text{earth}|\text{hubble}, \text{telescope}, \text{achievements})$



Estimation: The Sampling Game

- The query and the documents are samples from an unknown relevant model R
- The sampling game
 - Relevant examples are missing
 - Let $q = q_1 \dots q_n$ - the only sample that we have from the R distribution
 - Unknown distribution P , sample 3 times, get: **hubble**, **telescope**, **achievements**
 - What's the chance we get a word **earth** if we sample one more time?
- $$P(\text{earth}|R) \approx \frac{p(\text{earth}|\text{hubble}, \text{telescope}, \text{achievements})}{p(\text{hubble}, \text{telescope}, \text{achievements})}$$



Estimating the Joint Distribution

- $p(w, q_1 \dots q_n) = ?$

Estimating the Joint Distribution

- $p(w, q_1 \dots q_n) = ?$
- Assume that the query words $q_1 \dots q_n$ and the words w in relevant documents are sampled identically and independently from a unigram distribution M_R
- Let M represent some finite universe of unigram distributions from which we could sample
- Sampling process: pick one distribution $M \in M$ with probability $p(M)$ and sample from it $n+1$ times
- $p(w, q_1 \dots q_n) = \sum_{M \in M} p(M)p(w, q_1 \dots q_n|M)$

I.I.D Sampling

- $p(w, q_1 \dots q_n) = \sum_{M \in \mathcal{M}} p(M) p(w, q_1 \dots q_n | M)$
- Using the I.I.D assumption we get that the conditional probability is

$$p(w, q_1 \dots q_n | M) = p(w | M) \prod_{i=1}^n p(q_i | M)$$

- Now we can estimate the joint probability
- $p(w, q_1 \dots q_n) = \sum_{M \in \mathcal{M}} p(M) p(w | M) \prod_{i=1}^n p(q_i | M)$

Shortcut for the Estimation

- $p(w, q_1 \dots q_n) = \sum_{d \in D_{init}} p(M_d) p(w|M_d) \prod_{i=1}^n p(q_i|M_d) =$
 $\sum_{d \in D_{init}} p(M_d) p(w|M_d) p(q|M_d) = \sum_{d \in D_{init}} p(w|M_d) p(M_d|q) p(q)$
- Recall that $p(w|R) \approx \frac{p(w, q_1, \dots, q_n)}{p(q)}$
- Therefore, $p(w|R) \approx \sum_{d \in D_{init}} p(w|M_d) p(M_d|q)$
- $p(M_d|q) = \frac{p(q|M_d) p(M_d)}{\sum_{d \in D_{init}} p(M_d) p(q|M_d)}$

Summary

- Basic relevance model: (RM1)
- $p(w|R) \approx \sum_{d \in D_{init}} p(w|M_d)p(M_d|q)$
- $p(M_d|q) = \frac{p(q|M_d)p(M_d)}{\sum_{d \in D_{init}} p(M_d)p(q|M_d)}$
- Interpolated relevance model (RM3, Abdul Jaleel et al. '04)
- $p(w|R) \approx \beta p_{MLE}(w|q) + (1 - \beta) \sum_{d \in D_{init}} p(w|M_d)p(M_d|q)$

Ranking


- Compute the KL divergence between a document language model and R
- $Score(d) = -D(p(\cdot | R) || p(\cdot | M_d))$
- $Score(d) = -\sum_w p(w|R) \ln \frac{p(w|R)}{p(w|M_d)}$

Empirical Performance

	ROBUST			WT10G		
	MAP	p@5	NDCG@5	MAP	p@5	NDCG@5
Initial	24.8	47.9	50.0	19.6	33.4	32.8
RM1 (OptAQ)	23.4	48.6	50.9	16.6 ⁱ	33.6	31.2
RM3 (OptAQ)	28.1 _r ⁱ	49.8	51.9	21.7 _r ⁱ	34.2	33.0
RM1 (OptPQ)	28.7 ⁱ	56.5 ⁱ	59.5 ⁱ	23.2 ⁱ	43.1 ⁱ	41.4 ⁱ
RM3 (OptPQ)	34.4 _r ⁱ	59.6 _r ⁱ	62.5 _r ⁱ	30.3 _r ⁱ	49.1 _r ⁱ	47.6 _r ⁱ
RM1 (LOO)	23.4	48.6	50.9	16.6 ⁱ	33.6	31.2
RM3 (LOO)	27.3 _r ⁱ	48.5	50.5	21.2 _r	33.8	32.2
RM1 (10Fold)	23.4	48.6	50.9	15.7 ⁱ	32.0	29.4
RM3 (10Fold)	27.9 _r ⁱ	49.2	51.3	20.4 _r	34.0	32.1

Table 1: OptAQ– optimizing MAP over all queries, OptPQ– optimizing MAP per query, LOO– leave-one-out using MAP for optimization, 10Fold– 10-fold cross validation. i,r - statistically significant differences with Initial and RM1.

Relevance Feedback



- Allows the system to update the original query representation
 - May be iterative
 - May be implicit (e.g., from clicks)
- Pseudo-relevance feedback:
 - No judgments given, but we assume relevance of the first documents retrieved

Relevance Feedback

- Pseudo feedback

- Retrieve an initial set of n documents
- Assume that the top m ranked documents are relevant
- Induce query model

- $$P(M_d|q) = \frac{P(q|M_d)P(M_d)}{\sum_{d \in D_{init}} P(M_d)P(q|M_d)}$$

- True feedback

- Retrieve an initial set of n documents
- The user marks m relevant documents
 - Top-down process
- Induce query model

- $$P(M_d|q) = \frac{1}{m}$$

“Relevant”

Rank	Doc
1	d1
2	d2
3	d3
4	d4
5	d5
6	d6

Relevant

Relevant

Relevant

Rank	Doc
1	d1
2	d2
3	d3
4	d4
5	d5
6	d6

Example

- Rank the following 4 documents in response to the query “orange apple”
 - d1: orange orange lemon lemon clementine apple
 - d2: orange and red are beautiful colors
 - d3: orange and cellcom offer a special deal
 - d4: the fruits I like most are orange and apple
 - Construct a relevance model using the top 2 documents

Relevance Model Example

- Rank the following 4 documents in response to the query “orange apple”
 - We assume: Dirichlet smoothing with $\mu=1000$

$$p_{Dir}(w|M_d) = \frac{tf(w \in d) + \mu p_{MLE}(w|M_C)}{dl_d + \mu}$$

- $p(q|d1) = p(orange|d1) * p(apple|d1) = \left(\frac{2 + \frac{5}{28}1000}{6 + 1000} \right) * \left(\frac{1 + \frac{2}{28}1000}{6 + 1000} \right) = 0.0129$
- $p(q|d2) = p(orange|d2) * p(apple|d2) = \left(\frac{1 + \frac{5}{28}1000}{6 + 1000} \right) * \left(\frac{0 + \frac{2}{28}1000}{6 + 1000} \right) = 0.01267$
- $p(q|d3) = p(orange|d3) * p(apple|d3) = \left(\frac{1 + \frac{5}{28}1000}{7 + 1000} \right) * \left(\frac{0 + \frac{2}{28}1000}{7 + 1000} \right) = 0.01264$
- $p(q|d4) = p(orange|d4) * p(apple|d4) = \left(\frac{1 + \frac{5}{28}1000}{9 + 1000} \right) * \left(\frac{1 + \frac{2}{28}1000}{9 + 1000} \right) = 0.01277$

Relevance Model Example

- Construct a relevance model using the 2 most highly ranked documents

$$P(w|R) \approx \sum_{d \in D_{init}} P(w|M_d)P(M_d|q)$$
$$P(M_d|q) = \frac{P(q|M_d)P(M_d)}{\sum_{d \in D_{init}} P(M_d)P(q|M_d)}$$

- $P(clementine|R) = \frac{1}{6} * \frac{0.0129}{0.0129+0.01277} = 0.0838$
 - $P(fruit|R) = \frac{1}{9} * \frac{0.01277}{0.0129+0.01277} = 0.055$
 - $P(apple|R) = \frac{1}{6} * \frac{0.0129}{0.0129+0.01277} + \frac{1}{9} * \frac{0.01277}{0.0129+0.01277} = 0.139$
 - ...
- Documents are ranked according to the similarity of their model to that of the query

Clipping the Query Model


- The expanded query contains many terms
 - The amount of unique terms in documents provided for feedback
- Problems:
 - Query drift – why?
 - Computationally demanding
- Clipping the query model
 - Use the top v terms assigned the highest probability
 - Underlying assumption: these terms are most likely to represent the information need
 - $p(w|RM3) \stackrel{\text{def}}{=} \beta p_{MLE}(w|q) + (1 - \beta)p(w|RM1^{clipped})$

Clipping - Example

- q_0 = “nobel prize”
- INEX collection
 - # of documents = 2,666,190
 - # of relevant documents for q_0 = 22
- Construct RM3 model
 - 3 relevant documents serve as true feedback
 - $\beta = 0.2$
 - Clipping: $v = 100$
- statistics:
 - 79 out of 100 terms are stopwords!
 - The top 12 terms with the highest probability are:
 - nobel, prize, the, of, in, and, for, on, 2007, or, medicine, physiology

Additional Thoughts

- Given examples of relevant documents provided by the user
 - Will the expanded query always be closer to the centroid of the relevant documents?
- Suggest an approach for ameliorating this drawback
- How would you estimate $P(M_d | q)$?
- Using KL for ranking results in an IDF effect while using the MLE for the query language model
 - Rank equivalent to the query likelihood approach
- Does the IDF effect hold for any query model?
 - Raiber and Kurland 2017



Does the IDF Effect Hold for Any Query Model?

- M_d : Smoothed unigram document language model
- $M_d \stackrel{\text{def}}{=} (1 - \lambda)M_d^{MLE} + \lambda M_C^{MLE}$
- For $w \in d$:
 - $p_s(w|M_d) = (1 - \lambda) p(w| M_d^{MLE}) + \lambda p(w| M_C^{MLE})$
- For $w \notin d$:
 - $p_u(w|M_d) = \lambda p(w| M_C^{MLE})$
- $-CE(M_q, M_d) = \sum_w p(w|M_q) \log p(w|M_d) =$
 $\sum_{w \in d} p(w|M_q) \log p_s(w|M_d) + \sum_{w \notin d} p(w|M_q) \log p_u(w|M_d) =$

Does the IDF Effect Hold for Any Query Model?

- $-CE(M_q, M_d) = \sum_w p(w|M_q) \log p(w|M_d) =$
 $\sum_{w \in d} p(w|M_q) \log p_s(w|M_d) + \sum_{w \notin d} p(w|M_q) \log p_u(w|M_d) =$
 $\sum_{w \in d} p(w|M_q) \log \frac{p_s(w|M_d)}{p_u(w|M_d)} + \sum_w p(w|M_q) \log p_u(w|M_d) =$
 $\sum_{w \in d} p(w|M_q) \log \frac{p_s(w|M_d)}{\lambda p(w|M_C^{MLE})} + \log \lambda - CE(M_q, M_C^{MLE}) \propto$
 $\sum_{w \in d} p(w|M_q) \log \left(1 + \frac{1 - \lambda}{\lambda} \frac{p(w|M_d^{MLE})}{p(w|M_C^{MLE})} \right) + \log \lambda$